

Concurrency Control Techniques

Sanjay Kumar Gupta

Assistant Professor - PSIT

1

Concurrency Control

When **several transaction** execute **simultaneously** then there is a **risk of violation of data integrity**.

- **Concurrency Control** is a **procedure of managing simultaneous transactions** ensuring their **Atomicity, Isolation, Consistency & Serializability**.
- In **absence of concurrency control mechanism**, we may face:

Concurrency Problems in Transactions

1. **Dirty Read Problem**
2. **Lost Update Problem**

Lock Based Protocol

A **Lock** is a mechanism to **control concurrent access** to a data item.

1. Shared Mode (S):

-> **Only read operation** on data is applicable

Transaction T1	
LOCK S (A)	// Shared Lock on data "A"
R (A)	// Read Operation on data "A"
Unlock (A)	// Unlock data "A"

2. Exclusive Mode (X):

-> **Both read & write operations** on data are applicable.

Transaction T1	
LOCK X (A)	// Exclusive Lock on data "A"
R (A)	// Read Operation on data "A"
W (A)	// Write Operation on data "A"
Unlock (A)	// Unlock data "A"

3

Lock Compatibility matrix

Suppose **T1** and **T2** are **parallel transactions**, and **both** want to perform read and write operations on the same data, say "A".

A shared lock is denoted by "**S**," and an Exclusive lock is denoted by "**X**".

		Request →	
Grant ↓		Shared (S)	Exclusive (X)
	Shared (S)	✓	✗
	Exclusive (X)	✗	✗

Compatibility Lock Table

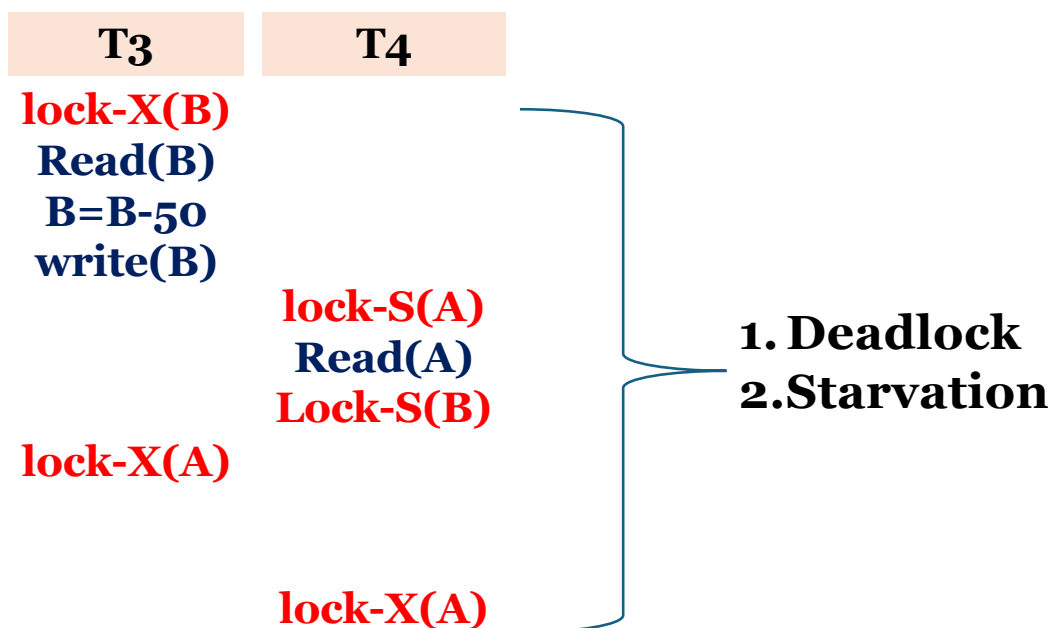
Examples of a Transaction Using Locks

T ₁
Lock-S(A)
Read(A)
Unlock(A)
Lock-S(B)
Read(B)
Unlock(B)
Display(A+B)

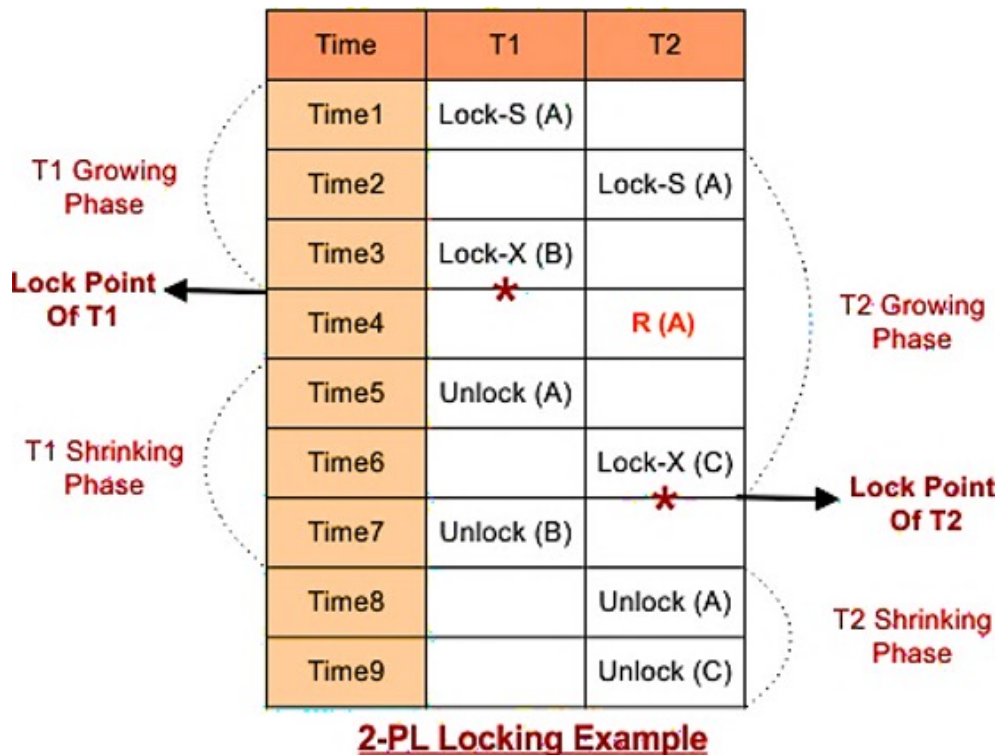
5

Drawbacks of Lock Based Protocol

Consider the given *partial* schedule,



Two-Phase Locking Protocol



7

Two-Phase Locking Protocol

- 2-PL is an **extension of Shared/Exclusive locking**.
- It is used to **reduce the problems** of Shared/Exclusive locking

Any schedule that follows **2PL** will **always** be **serializable**, which **was not** in Shared/Exclusive locking

Phases of 2PL

- I. Growing Phase:** In the Growing phase, **only Locks are acquired** by a transaction and **no locks are released** by a transaction at that time.
- II. Shrinking Phase:** In the shrinking phase, **only Locks are released by transaction**, and **no locks are acquired** by a transaction at that time.

Rigorous 2-PL

It is like **Strict 2-PL** in its advantages and disadvantages **but** a little bit **more strict** than Strict 2-PL.

- It must satisfy the **Basic 2-PL**
- Each transaction **should hold all Exclusive(X) Locks** as well as **Shared(S) Locks** until the Transaction is **Commits** or **Aborted**.

T1	T2
Lock-S(A)	
READ (A)	
	Lock-S(A) // Wait
LOCK-X(B)	
READ (B)	
WRITE (B)	
Commit (A)	
Commit (B)	
Unlock (A) ▼	
Unlock (B)	Lock-S(A) // Granted ▼
	READ (A)
	Commit (A)
	Unlock (A)

9

Strict 2-PL

A schedule will be in **Strict 2PL** if,

- It must satisfy the **Basic 2-PL**
- Each transaction should hold all **Exclusive(X) Locks** until the Transaction is **Committed** or **Aborted**.


T1	T2
Lock-S(A)	
READ (A)	
	Lock-S(A) // Granted
LOCK-X(B)	
READ (B)	
WRITE (B)	
Commit (A)	
Commit (B)	
Unlock (A)	
Unlock (B)	

Conservative 2-PL

The idea is there is **no growing phase**, transaction **start** directly from **lock point**, i.e., **transaction must first acquire all the required locks then only it can start execution.**

Shrinking phase will work as usual, and **transaction can unlock any data item anytime.**

We must have a knowledge in future to understand what is data required so that we can use it.



Other Protocols

Lock based protocols **ensure Conflict Serializability**, but it causes two problems:

1. **Deadlock** &
2. **Starvation**

The **alternative approaches** to control **concurrency** are:

- A. Timestamp-Based Protocols**
- B. Validation Based Protocols**

13

Time Stamping Protocols

- ❖ The **Timestamp Ordering Protocol** arranges the transactions based on their respective **timestamps**.
- ❖ Basic idea of time stamping is **to decide the order** between the transaction before they enter in the system using a stamp (time stamp), **in case of any conflict during the execution order can be decided using the time stamp**.

Example: Suppose three transactions, T₁, T₂, and T₃, execute in parallel on the same data A.

- “T₁ arrived at 8:00 AM. Please assign a timestamp value of 100 and call it ‘**Older**.’
- “T₂ arrived at 8:10 AM, so please assign a timestamp value of 200 and name it ‘**Younger**.’
- “T₃ arrived at 8:15 AM. Therefore, assign a timestamp value of 300 and name it ‘**Youngest**.’

Time Stamping Protocols

- ❖ Let's understand how this protocol works, here we have two idea of timestamping, **one for the transaction**, and **other for the data item**.

Time stamp for transaction

- With each transaction t_i in the system, we associate a unique fixed timestamp \rightarrow denoted by $TS(t_i)$.
- If a transaction has been assigned a timestamp $TS(t_i)$ and a new transaction t_j , enters the system with a timestamp $TS(t_j)$, then always $TS(t_i) < TS(t_j)$

15

Time Stamping Protocols

Two things are to be noted:

1. Time stamp of a transaction **remain fixed** throughout the execution
2. It is **unique** means **no two transaction can have the same timestamp**

Time stamp with data item

- **W-timestamp(Q)** is the **largest time-stamp** of any transaction that executed $write(Q)$ successfully
- **R-timestamp(Q)** is the **largest time-stamp** of any transaction that executed $read(Q)$ successfully

These timestamps are **updated** whenever a **new read(Q)** or **write(Q)** instruction is executed.

Time Stamping Protocols

Suppose a transaction T_i request a *read(Q)*,

- If $TS(T_i) < W\text{-timestamp}(Q)$, then T_i needs to read a value of Q that was already overwritten.
Hence, the read operation is rejected, and T_i is rolled back
- If $TS(T_i) \geq W\text{-timestamp}(Q)$, then the read operation is executed, and $R\text{-timestamp}(Q)$ is set to the **maximum** of $R\text{-timestamp}(Q)$ and $TS(T_i)$.

17

Time Stamping Protocols

Suppose a transaction T_i request a *write(Q)*,

- 1: If $TS(T_i) < R\text{-timestamp}(Q)$
Write rejected; T_i rolled back (value needed earlier)
- 2: If $TS(T_i) < W\text{-timestamp}(Q)$
Write rejected; T_i rolled back (obsolete value)
- 3: If $TS(T_i) \geq R\text{-timestamp}(Q)$
Write operation is executed,
and **$W\text{-timestamp}(Q)$ is set to $TS(T_i)$**
- 4: If $TS(T_i) \geq W\text{-timestamp}(Q)$
Write operation is executed,
and **$W\text{-timestamp}(Q)$ is set to $TS(T_i)$**

Thomas Write Rule

Improvement over the timestamp-ordering protocol.

Increases **concurrency** and may generate **view-serializable** schedules.

Key Features:

- **Blind writes** can be ignored under specific conditions.
- **Read rules** remain unchanged.
- **Write rules** slightly modified from timestamp-ordering protocol.

When T_i attempts to write data item Q ,

- if $TS(T_i) < W\text{-timestamp}(Q)$, then T_i is attempting to write an obsolete value of $\{Q\}$.

Rather than *rolling back T_i as the timestamp ordering protocol would have done*, **this {write} operation can be ignored.**

19

Timestamp Ordering Protocol

Advantages of Timestamp Protocol

- The Timestamp protocol **ensures serializability** and **Deadlock removal** because a transaction with a smaller timestamp (TS) **executes before a transaction with a higher TS.**

Disadvantages of Timestamp Protocol

- The schedule **may not be recoverable.**
- The schedule **may not be cascading-free.**

Validation Based Protocol

It is also known as **optimistic concurrency control technique**.

- Used where transactions are **READ** only & **CONFLICTS** are low.

❖ Every T_i must go through three phases:

1. Read & Execution Phase

- Transaction T_i writes only to the **temporary local variable**.
 - It means, it does not update actual database.

2. Validation Phase

- T_i performs **validation test** to determine if **local variables can be written without violating serializability**.

3. Write Phase

- If **T_i is validated**, the updates are **applied to the database**,
 - or **T_i is rolled back**.

21

Validation Based Protocol

Each Transaction has **three** timestamps:

1. **start (T_i)**: the time when T_i started its execution
2. **validation (T_i)**: the time when T_i entered validation phase
3. **finish (T_i)**: the time when T_i finished its write phase

To clear the **validation test** by T_i , must follow **anyone** conditions:

- (i) **Finish (T_i) < Start (T_j)**: It means T_i is Older transaction & It get finished before T_j starts. **(NO OVERLAP)**
- (ii) **Finish (T_i) < Validate (T_j)**: This ensures actual write by T_i & T_j will **NOT OVERLAP**.
- (iii) **Validate (T_i) < Validate (T_j)**: It ensures that T_i has completed Read phase before T_j completes Read Phase.

Validation Test

T_{14}	T_{15}
read(B)	read(B)
	$B := B - 50$
	read(A)
	$A := A + 50$
read(A)	
$\langle \text{validate} \rangle$	$\langle \text{validate} \rangle$
display(A + B)	write(B)
	write(A)

ADVANTAGES

- > Maintains Serializability
- > Free from Cascade Rollback
- > Less Overhead than other Protocols

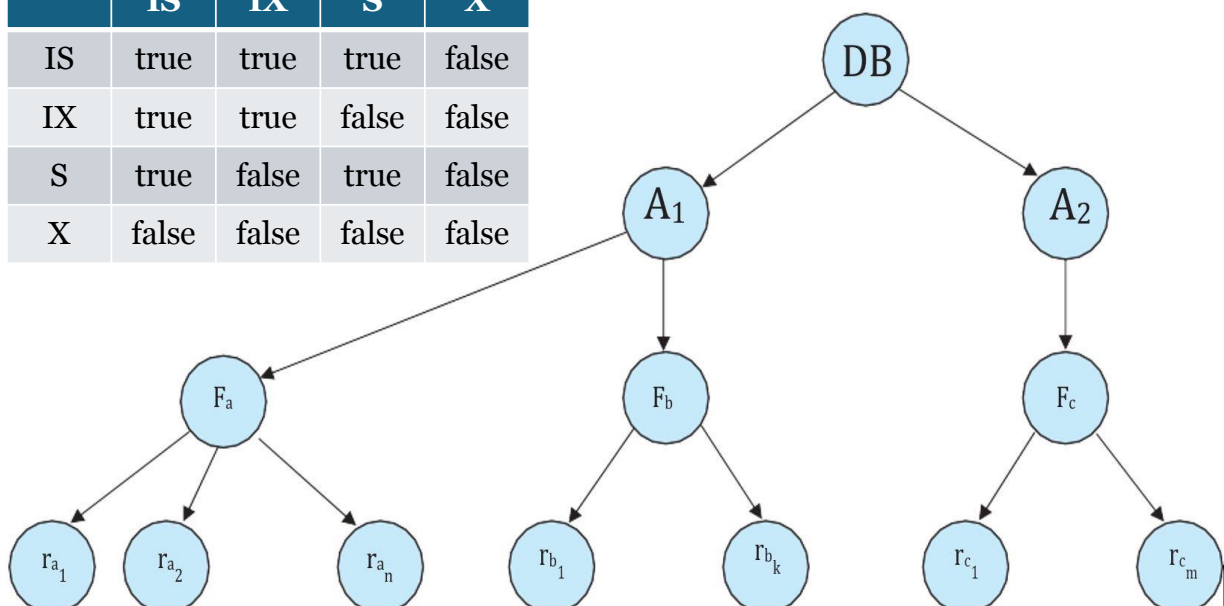
DISADVANTAGES

- > Starvation of long transaction due to conflicting transaction

23

Multiple Granularity

	IS	IX	S	X
IS	true	true	true	false
IX	true	true	false	false
S	true	false	true	false
X	false	false	false	false



Multiple Granularity

Granularity: It is the size of data item allowed to lock.

- It can be defined as hierarchically breaking up the database into blocks which can be locked.
- The Multiple Granularity protocol enhances concurrency and reduces lock overhead.
- It maintains the track of what to lock and how to lock.
- It makes easy to decide either to lock a data item or to unlock a data item.
- This type of hierarchy can be graphically represented as a tree.